

The pivot: scored offline, traced online

the pivot

Same bot, same code. The pivot is the QUESTION, not the tool:
'did it get this right' (needs an answer key) becomes
'what happened, and what did it cost' (needs an instrument).

Part 2 builds that instrument: nested spans with token, cost, and
latency accounting, and the same eval logic watching for drift.

A trace does not replace a test set -- it answers what a test set structurally cannot.