

Offline metrics: exact match, recall, and the BLEU trap

three offline metrics, three different questions

exact match	did the prediction equal the expected string, verbatim
recall@k	of the truly relevant items, how many did retrieval surface
BLEU / ROUGE	surface n-gram overlap with a reference -- not meaning

All three are pure functions of (prediction, expected). No model is in the scoring path.