

Offline metrics: exact match, recall, and the BLEU trap

why surface metrics mislead

BLEU/ROUGE grade word choice, not correctness.

A reworded, equally correct answer can score near zero.

A wrong answer that happens to reuse the reference's words can score high.

Measure the thing you actually care about, not its surface form.

This is the deterministic version of a well-known failure mode -- no judge, no LLM, just token sets.