

Hosted vs local

two ways to run inference

hosted API (OpenAI, Anthropic, a cloud endpoint):

a network call to someone else's GPUs and ops team.

local model (this deck's gpt2, or an Ollama model):

a process call to weights you loaded onto your own machine.

Every prior topic in this deck used the local path: gpt2, in-process, on CPU.