

litellm, one shape

one call shape, any provider

```
completion(  
    model="ollama/qwen2.5:0.5b",  
    messages=[{"role": "user", "content": ...}],  
    temperature=0,  
    api_base="http://localhost:11434",  
)  
  
swap model="ollama/..." for "openai/gpt-4o" or "anthropic/claude-..."  
and this exact call hits a hosted provider instead.
```

litellm normalizes every provider behind the OpenAI chat-completion shape.