

Quantization

the real gpt2 parameter count

```
n_params = sum(p.numel() for p in model.parameters())  
n_params = 124,439,808
```

the same count either way -- quantization changes how each number is STORED, never how many numbers there are.

`model.parameters()` walks the real loaded `nn.Module`, not a spec sheet.