

The KV cache

what gets cached

```
past_key_values: one (key, value) pair per layer x 12 layers  
each holds every token's key/value the model has already seen
```

```
step 1 (the prefix): no cache yet, so it computes all positions.
```

context: "The cat sat on the", 5 tokens, 12 layers to fill.