

vLLM, the throughput path

the outgrow path

vLLM: an inference server built around PagedAttention + continuous batching, purpose-built to push many concurrent requests through one GPU at high utilization.

same OpenAI-shaped call surface litellm already showed --
point api_base at a vLLM server instead of ollama, nothing else changes.

The serving engine changes underneath; the call shape from litellm-one-shape does not.