

Why vLLM is fast

the naive server

naive serving: one contiguous KV-cache allocation per request,
sized for the longest sequence the server might see.

most of that block sits empty for most requests.
a batch only starts once, and only ends once, together.

Fixed-size, over-provisioned, and static-batched: memory and requests both wait on the slowest member.