Why vLLM is fast

vLLM's two ideas

```
PagedAttention: the KV cache is sliced into fixed-size BLOCKS,
allocated on demand, like OS virtual-memory pages.

continuous batching: a finished request's slot is handed to the
next queued request immediately, not at the next batch boundary.
```

Both ideas raise GPU utilization by shrinking wasted memory and wasted idle time, not by changing the math.