

Why vLLM is fast

what does NOT change

```
same forward pass, same logits, same greedy/sampled token  
vLLM changes THROUGHPUT (requests/sec, GPU $ per token),  
never the model's answer for a given prompt and settings.
```

This deck's oracles never assert a wall-clock number for vLLM: no vLLM install here to measure honestly.